Stock Market Trading with LSTM Analysis

Franklin Doane
dept. of Computer Science
Tennessee Technological University
Cookeville, United States
fgdoane42@tntech.edu

I. DATASET CREATION

A. Generate stock list

1) Download stocks by market cap: I downloaded a list of all stocks ranked by market cap off the internet, and saved it as a .tsv

2) Read tsv: Run the stock list builder, which starts by reading in the file as a pandas dataframe

3) Iterate while stock list is < 1500: Create empty list for stock symbols. Continue to perform the following while the stock list len is < 1500.

4) Grab next highest symbol: Grab next symbol from ranking dataframe.

5) Download data: Attempt to download day candles going back to the date 01/01/2011 - Δ(days=100).

6) Check dates captured: Check if the download captured 01/01/2011 (oldest date to be used) and 12/31/2024 (most recent date to be used).

7) Check context captured: Check that there are at least 50 days of data before the first date (so that the first training example will be for 01/01/2011)

8) Judge stock: If stock data download captured data range and addition context, it gets added to the list. Otherwise it gets rejected.

9) Save list: Once 1500 stocks have been accepted into the list, save it as a text file of tickers.

B. Setup downloads

1) Decide on train set date ranges: I made a .md file that numbers and outlines the date range for each dataset.

2) Decide on eval set date ranges: Each evaluation set is 1 year long and has a first date that is 1 year after the end of the training set data is corresponds with.

3) Create bash files: I created 2 bash files for training and evaluation that run the download script for each dataset with it's unique arguments.

C. Download dataset (for each training and eval set)

1) Collect args: Collect arguments pertaining to date range and save location.

2) Load stock list: Load stock list text file as a list.

3) Open thread pool: Open pool of threads for calling downloads like follows based on each stock in the list.

4) Call download: Download from (01/01/start_year - context*2), to 12/31/end_year with auto-adjust = false

5) Update indexing structure: Reset columns to be base OHLCV. Reset df index.

6) Get important data: Index needed data out of dataframe with (first day of start_year - context) : (last day of end_year)

7) Save: Save stock data to csv file

II. TRAINING SETUP

A. Load Config

1) Load: Load config as json object with path from positional arg

2) Unpack: Unpack config keys into variables

B. Setup regulator

1) Init regulator: Pass starting learning rate and other kwargs from config

C. Load dataset

1) Init dataset object: The dataset object is initialized and given config. Then the load dataset method is called which does the following.

2) Verify files: Loads all csvs in dataset and verifies a minimum length. A legacy feature at this point.

3) Split files into task groups: List of valid files is split into groups to have each group loaded with threads.

4) Open dataset thread pool: Open thread pool to load files.

5) Load csv files: Open every file from directory listed in config. Load as pandas dataframe with date index column and put them all in a list.

6) Calculate indicators: Calculate all additional features listed in configs. Cut out any data at the front of the dataframes that is missing these features (e.g. for rolling avgs).

7) Create date feature: 1hot encode day of the week and add the new weekday features to all dataframes.

8) Slice data into training examples: Use rolling window to cut each dataframe into training example dataframes with n day history. All examples get added to a list and labels are made into a dataframe of their own.

9) Normalize examples: Each dataframe training example is normalized using norm function from config.

10) Create tensor input: Each dataframe is converted to numpy then a tensor with the datatype and device listed in the config. All labels are converted a single tensor as well.

D. Load model

1) Dynamically load model class: The Python file containing the model class is loaded by name using importlib.

2) Call init: Init custom model class. The custom model class is passed a name, starting LR, and input feature count from dataset. The model inherits from torch.nn.Module which is also inited.

3) Setup pytorch objects: Model class init sets up LSTM object and Sequential object for forward passes, a default optimizer, a null loss func and an empty history.

4) Set optimizer: Call model set optim method to initialize and save optimizer to model class.

5) Send to device: Set model to device specified in config.

E. Setup scheduler

1) Init scheduler: Pass optimizer from model. Give it kwargs from config

F. Setup dataloader

1) Check for test set: Check's to see if the model save config contains test set indices. For loading a partially trained model. Indices get set after init if none are there. Pretty much a vestigial feature.

2) Call init: Init object and pass in dataset object (and test set indices if applicable)

3) Pick test indices: If not test set indices were passed, create list of indices and either random sample indices (legacy) or slice chunk off the front to create test set indices.

4) Seperate test set: Use indices to make a mask and remove test examples from all indices to get training set indices

5) Shuffle train set: random.shuffle train set indices

6) Split batches: Split train set indices into batches using a list comprehension and slicing with option to drop last partial batch. Transforms list[int] -¿ list[list[int]]. Casts result to tensor.

7) Setup first inputs: Create empty list for inputs and labels. Prefetch first n batches by indexing them out of dataset and adding them to input and label sets. Remove these indices from batch indices.

8) Load test set: Index test set and labels out of data and Store

III. TRAINING LOOP (PER EPOCH)

A. Setup

1) Get regulator coefs: Pass regulator lr and get the coefficient for each logit. Coefficients will be stepped down if LR has changed

2) Update loss: Call function to init custom MSE loss and pass it the regularization coefficients. Call method to set new loss for model class.

B. Train batches

1) Unpack dataloader: Each iteration grabs a batch from the loaded set of batches. If the loaded set is empty, the dataloader loads indexes more batches out of the dataset. Continues until all batch indices are popped.

2) Forward pass: Call model forward with batch inputs

3) Call backward pass: Pass model backpass method the logits and batch labels (with added time series dimension). Starts by putting model in train mode.

4) Calc loss: Zero optimizer gradient then run model's built-in loss function. Backwards' the loss.

5) Optimize: Step optimizer.

C. Save

1) Checkpoint model: Save epoch copy of model as pkl to model save dir. Includes model params as well as model training history.

D. Validation

1) Eval setup: Set model in eval mode. Create data structures for saving outcomes and losses.

2) Open each item in test set: Iterate over test set inputs and labels. Do the following for each one

3) Forward pass: Call model on test input

4) Calculate loss: Grab close price (3rd label item), and pass high and low price into loss function

5) Update outcome counts: Decipher outcome as either win, stop-loss, both (loss), neither, missed win, missed loss, missed stop-loss, missed both, or missed neither. Also update the count for the broader category of loss, and update the counts for why the win was missed (predicted stop-loss, predicted neither, predicted both).

6) Final calculations: After all examples are run through, calculate the following statistics: win rate, random win rate, guess rate, average timeout, and average eval loss

7) Update scheduler: Send average example loss to scheduler.

IV. EVALUATION SETUP

A. Load configs

1) Load eval config: Load config as JSON object.

2) Unpack eval config: Unpack values from config as variables.

3) Load training config: Load the config that was used for training the model as a JSON object.

B. Load model

1) Dynamically load model class: The Python file containing the model class is loaded by name using importlib.

2) Call init: Init custom model class. The custom model class is passed a name, starting LR of 0, and input feature count from config. The model inherits from torch.nn.Module which is also inited.

3) Load model save: Loads model params from pkl file.

C. Start agent

1) Call agent init: Initialize agent object and pass config, model, and starting balance.

2) Unpack agent parameters: Agent init unpacks all of it's rules for buying and selling from config.

3) Create data structures: Agent init creates empty data structures for tracking outcomes, balance, owned stocks, and labels for owned stocks.

D. Load dataset

1) Init dataset and load: The dataset object is initialized and given config. The load method is called which does the following.

2) Read csvs: The dataset object reads all csvs in dataset directory and store them in a dict with the format {"ticker": dataframe}

3) Calculate additional features: Iterate through dataset keys and calc all additional features specified in training config.

4) 1-hot encode dates: For each dataframe in dataset, keep date column as index but add 1-hot encoded day of the week columns.

5) Setup comparison stock: Look for saved file of balances from investing with comparison stock (e.g. S&P). If not, load comparison stock data from dataset.

E. Setup for iteration

1) Init date range object: Initialize custom date range iterator and pass it the start and end dates to iterate over. Each date in the iteration will be the day to predict the outcome of.

V. EVALUATION LOOP (PER DAY IN DATE RANGE)

A. Get inputs

1) Attempt to index date: For each stock dataframe, attempt to index the current date in the simulation. If it fails, return no data (common for weekends or holidays).

2) Get input sequence: If date was indexed successfully, grab 50 day sequence prior for model input.

3) Get last price: Index the close price of last day of sequence. This will be considered the purchase price if stock is purchased.

4) Normalize input sequence: Use normalization function from training to normalize sequence.

5) Turn input into tensor: Drop date string index column from dataframe and convert to torch tensor.

6) Get label data: Index current date row out of dataframe and turn into dict to use for labels.

7) Compile data: Put all inputs and label from all stocks for the day into an iterable.

B. Iterate through inputs (for each stock)

1) Unpack data: Unpack stock ticker, input, final price in input sequence, and day candlestick (labels).

2) Forward pass: Forward pass input through model.

3) Save data: Save model logits, and all other unpacked stock data into data structures.

C. Purchase

1) Call agent buy: Tell agent to make purchases. Pass model predictions and labels.

2) Filter predictions: The agent buy function filters predictions for ones that meet the criteria for purchasing.

3) Allocate funds: If there are stocks that meet the criteria, divide current amongst them equally. Agent will remove stocks if there are too many to have minimum $1 per purchase.

4) Finalize purchases: Add purchased stocks and their purchase amount to agent stocks. Save label data as well.

D. Sell

1) Call agent sell: Tell agent to sell stocks.

2) Calculate aim prices: Calculate prices that agent aims to sell stocks at, as well as prices it will stop loss at.

3) Calculate increases: Use label data to calculate what the increase (positive or negative) would've been for each stock.

4) Calculate sale totals: Use buy totals and increases to calculate sale totals. Sum them to get new balance.

5) Update outcome counts: Update agent's stored counts of different outcome types for statistical analysis.

6) Update increase averages: Update agent's stored averages of increase percentages for different outcome types.

7) Remove agent stocks: Remove all stocks and labels from agent's owned list.

E. Update

1) Update agent history: Add new balance, buy count, and increases to history for charting.

2) Update comparison stock: If generating data to compare to agent, calculate what the stocks increase would've been and update balance.

VI. EVALUATION AFTERMATH

A. Files

1) Calculate win rate: Use outcome counts to calculate win rate.

2) Save all data: Save balances, buy counts, increases, outcome counts, and generated comparison data to results directory.

B. Charts

1) Create and save charts: Generate chart and save to results directory.